

EDA_Complete

November 29, 2025

1 Notebook 1: Exploratory Data Analysis (EDA) and Data Preparation

1.1 MIMIC-III Clinical Database for Recommendation Systems

Author: Data Science Team

Date: November 2025

Objective: Complete EDA with data loading, inspection, cleaning, and train-test split

1.2 Workflow Overview

EDA WORKFLOW PIPELINE

1. DATA LOADING
↓
2. INITIAL INSPECTION (`head`, `info`, `describe`)
↓
3. MISSING VALUES & DUPLICATES
↓
4. OUTLIERS DETECTION
↓
5. EXPLORATORY ANALYSIS (Distributions)
↓
6. CORRELATION ANALYSIS
↓
7. CLASS BALANCING ASSESSMENT
↓
8. DATA CLEANING & PREPROCESSING
↓
9. TRAIN-TEST SPLIT & EXPORT
↓
10. FINAL SUMMARY REPORT

1.3 Section 1: Import Required Libraries

```
[30]: # Import essential libraries for EDA and data preprocessing
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
import os
import pickle
import json
from datetime import datetime
from scipy import stats
from scipy.stats import skew, kurtosis

# Machine Learning libraries
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from imblearn.over_sampling import SMOTE

# Configure visualization
warnings.filterwarnings('ignore')
sns.set_style('whitegrid')
plt.rcParams['figure.figsize'] = (12, 6)
plt.rcParams['font.size'] = 10

print(' All libraries imported successfully!')
print(f'Execution timestamp: {datetime.now().strftime("%Y-%m-%d %H:%M:%S")}' )
```

All libraries imported successfully!
Execution timestamp: 2025-11-27 16:46:38

1.4 Section 2: Load Dataset from MIMIC-III

```
[31]: # Load MIMIC-III clinical database tables
import kagglehub
from kagglehub import KaggleDatasetAdapter

DATASET_NAME = 'ihssanened/mimic-iii-clinical-databaseopen-access'

def load_mimic_table(file_name, date_col=None):
    """Load a single MIMIC-III table using KaggleHub."""
    try:
        print(f' Loading: {file_name}...', end=' ')
        df = kagglehub.load_dataset(
            KaggleDatasetAdapter.PANDAS,
            DATASET_NAME,
            file_name
```

```

    )
    if date_col and date_col in df.columns:
        df[date_col] = pd.to_datetime(df[date_col], errors='coerce')
    print(f' {df.shape[0]} rows x {df.shape[1]} columns')
    return df
except Exception as e:
    print(f' Error loading {file_name}')
    return None

print('\n==== Loading MIMIC-III Tables ===\n')
df_admissions = load_mimic_table('admissions.csv', date_col='admittime')
df_patients = load_mimic_table('patients.csv')
df_labevents = load_mimic_table('labevents.csv')
df_d_labitems = load_mimic_table('d_labitems.csv')

if df_admissions is None or df_patients is None:
    print('\nCritical tables failed to load!')
else:
    print('\n All critical tables loaded!')

```

==== Loading MIMIC-III Tables ===

```

Loading: admissions.csv... 129 rows x 19 columns
Loading: patients.csv... 129 rows x 19 columns
Loading: patients.csv... 100 rows x 8 columns
Loading: labevents.csv... 100 rows x 8 columns
Loading: labevents.csv... 76,074 rows x 9 columns
Loading: d_labitems.csv... 76,074 rows x 9 columns
Loading: d_labitems.csv... 753 rows x 6 columns

```

```

All critical tables loaded!
753 rows x 6 columns

```

All critical tables loaded!

1.5 Section 3: Initial Data Inspection (head, info, describe)

```
[32]: # Perform initial inspection of key tables
def inspect_table(df, name):
    print(f'\n{"*80}')
    print(f'TABLE: {name.upper()}')
    print(f'{"*80}')
    print(f'\nShape: {df.shape[0]} rows x {df.shape[1]} columns')
    print(f'\nFirst 5 rows:')
    display(df.head())
    print(f'\nData types and missing values:')



```

```

print(df.info())
print(f'\nDescriptive statistics:')
display(df.describe(include='all').round(2))

if df_admissions is not None:
    inspect_table(df_admissions, 'admissions')
if df_patients is not None:
    inspect_table(df_patients, 'patients')

```

=====

TABLE: ADMISSIONS

=====

Shape: 129 rows × 19 columns

First 5 rows:

	row_id	subject_id	hadm_id	admittime	dischtime	\
0	12258	10006	142345	2164-10-23 21:09:00	2164-11-01 17:15:00	
1	12263	10011	105331	2126-08-14 22:32:00	2126-08-28 18:59:00	
2	12265	10013	165520	2125-10-04 23:36:00	2125-10-07 15:13:00	
3	12269	10017	199207	2149-05-26 17:19:00	2149-06-03 18:42:00	
4	12270	10019	177759	2163-05-14 20:43:00	2163-05-15 12:00:00	

	deathtime	admission_type	admission_location	\
0	NaN	EMERGENCY	EMERGENCY ROOM ADMIT	
1	2126-08-28 18:59:00	EMERGENCY	TRANSFER FROM HOSP/EXTRAM	
2	2125-10-07 15:13:00	EMERGENCY	TRANSFER FROM HOSP/EXTRAM	
3	NaN	EMERGENCY	EMERGENCY ROOM ADMIT	
4	2163-05-15 12:00:00	EMERGENCY	TRANSFER FROM HOSP/EXTRAM	

	discharge_location	insurance	language	religion	marital_status	\
0	HOME HEALTH CARE	Medicare	NaN	CATHOLIC	SEPARATED	
1	DEAD/EXPIRED	Private	NaN	CATHOLIC	SINGLE	
2	DEAD/EXPIRED	Medicare	NaN	CATHOLIC	NaN	
3	SNF	Medicare	NaN	CATHOLIC	DIVORCED	
4	DEAD/EXPIRED	Medicare	NaN	CATHOLIC	DIVORCED	

	ethnicity	edregtime	edouttime	\
0	BLACK/AFRICAN AMERICAN	2164-10-23 16:43:00	2164-10-23 23:00:00	
1	UNKNOWN/NOT SPECIFIED		NaN	
2	UNKNOWN/NOT SPECIFIED		NaN	
3	WHITE	2149-05-26 12:08:00	2149-05-26 19:45:00	
4	WHITE		NaN	

	diagnosis	hospital_expire_flag	has_charevents_data	
0	SEPSIS	0	1	

1	HEPATITIS B	1	1
2	SEPSIS	1	1
3	HUMERAL FRACTURE	0	1
4	ALCOHOLIC HEPATITIS	1	1

Data types and missing values:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 129 entries, 0 to 128
Data columns (total 19 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   row_id            129 non-null    int64  
 1   subject_id        129 non-null    int64  
 2   hadm_id           129 non-null    int64  
 3   admittime         129 non-null    datetime64[ns]
 4   dischtime         129 non-null    object  
 5   deathtime         40 non-null     object  
 6   admission_type   129 non-null    object  
 7   admission_location 129 non-null    object  
 8   discharge_location 129 non-null    object  
 9   insurance         129 non-null    object  
 10  language          81 non-null     object  
 11  religion          128 non-null    object  
 12  marital_status   113 non-null    object  
 13  ethnicity         129 non-null    object  
 14  edregtime         92 non-null     object  
 15  edouttime         92 non-null     object  
 16  diagnosis         129 non-null    object  
 17  hospital_expire_flag 129 non-null    int64  
 18  has_chartevents_data 129 non-null    int64  
dtypes: datetime64[ns](1), int64(5), object(13)
memory usage: 19.3+ KB
```

None

Descriptive statistics:

	row_id	subject_id	hadm_id	admittime	\
count	129.00	129.00	129.00		129
unique	NaN	NaN	NaN		NaN
top	NaN	NaN	NaN		NaN
freq	NaN	NaN	NaN		NaN
mean	28036.44	28010.41	152343.44	2154-04-08 23:17:16.744185856	
min	12258.00	10006.00	100375.00		2102-08-29 07:15:00
25%	12339.00	10088.00	128293.00		2128-11-04 16:05:00
50%	39869.00	40310.00	157235.00		2150-08-22 17:33:00
75%	40463.00	42135.00	174739.00		2180-03-15 22:35:00
max	41092.00	44228.00	199395.00		2202-10-03 01:45:00
std	14036.55	16048.50	27858.79		NaN

	dischtime	deathtime	admission_type	\
count	129	40	129	
unique	129	40	3	
top	2164-11-01 17:15:00	2126-08-28 18:59:00	EMERGENCY	
freq	1	1	119	
mean	NaN	NaN	NaN	
min	NaN	NaN	NaN	
25%	NaN	NaN	NaN	
50%	NaN	NaN	NaN	
75%	NaN	NaN	NaN	
max	NaN	NaN	NaN	
std	NaN	NaN	NaN	
	admission_location	discharge_location	insurance	language
count	129	129	129	81
unique	5	10	4	5
top	EMERGENCY ROOM	ADMIT	DEAD/EXPIRED	Medicare
freq	81	40	98	58
mean	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN
std	NaN	NaN	NaN	NaN
	religion			
count				128
unique				10
top				CATHOLIC
freq				59
mean				NaN
min				NaN
25%				NaN
50%				NaN
75%				NaN
max				NaN
std				NaN
	marital_status	ethnicity	edregtime	edouttime
count	113	129	92	92
unique	6	9	92	92
top	MARRIED	WHITE	2164-10-23 16:43:00	2164-10-23 23:00:00
freq	60	86	1	1
mean	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN
std	NaN	NaN	NaN	NaN
	diagnosis	hospital_expire_flag	has_charevents_data	
count	129	129.00	129.00	
unique	95	NaN	NaN	
top	SEPSIS	NaN	NaN	
freq	10	NaN	NaN	
mean	NaN	0.31	0.99	
min	NaN	0.00	0.00	
25%	NaN	0.00	1.00	

50%	NaN	0.00	1.00
75%	NaN	1.00	1.00
max	NaN	1.00	1.00
std	NaN	0.46	0.09

=====
TABLE: PATIENTS
=====

Shape: 100 rows × 8 columns

First 5 rows:

	row_id	subject_id	gender	dob	dod	\
0	9467	10006	F	2094-03-05 00:00:00	2165-08-12 00:00:00	
1	9472	10011	F	2090-06-05 00:00:00	2126-08-28 00:00:00	
2	9474	10013	F	2038-09-03 00:00:00	2125-10-07 00:00:00	
3	9478	10017	F	2075-09-21 00:00:00	2152-09-12 00:00:00	
4	9479	10019	M	2114-06-20 00:00:00	2163-05-15 00:00:00	

	dod_hosp	dod_ssn	expire_flag
0	2165-08-12 00:00:00	2165-08-12 00:00:00	1
1	2126-08-28 00:00:00	NaN	1
2	2125-10-07 00:00:00	2125-10-07 00:00:00	1
3	NaN	2152-09-12 00:00:00	1
4	2163-05-15 00:00:00	2163-05-15 00:00:00	1

Data types and missing values:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 100 entries, 0 to 99

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	row_id	100 non-null	int64
1	subject_id	100 non-null	int64
2	gender	100 non-null	object
3	dob	100 non-null	object
4	dod	100 non-null	object
5	dod_hosp	70 non-null	object
6	dod_ssn	77 non-null	object
7	expire_flag	100 non-null	int64

dtypes: int64(3), object(5)

memory usage: 6.4+ KB

None

Descriptive statistics:

row_id	subject_id	gender	dob	dod	\
--------	------------	--------	-----	-----	---

count	100.00	100.00	100	100	100
unique	NaN	NaN	2	99	100
top	NaN	NaN	F	2058-04-23 00:00:00	2165-08-12 00:00:00
freq	NaN	NaN	55	2	1
mean	20452.58	26162.33	NaN	NaN	NaN
std	10982.04	16201.83	NaN	NaN	NaN
min	9467.00	10006.00	NaN	NaN	NaN
25%	9526.50	10068.50	NaN	NaN	NaN
50%	20209.00	25128.00	NaN	NaN	NaN
75%	31391.25	42276.50	NaN	NaN	NaN
max	31872.00	44228.00	NaN	NaN	NaN
		dod_hosp	dod_ssn	expire_flag	
count		70	77	100.0	
unique		70	77	NaN	
top		2165-08-12 00:00:00	2165-08-12 00:00:00	NaN	
freq		1	1	NaN	
mean		NaN	NaN	1.0	
std		NaN	NaN	0.0	
min		NaN	NaN	1.0	
25%		NaN	NaN	1.0	
50%		NaN	NaN	1.0	
75%		NaN	NaN	1.0	
max		NaN	NaN	1.0	

1.6 Section 4: Missing Values and Duplicates Detection

```
[33]: # Analyze missing values
def analyze_missing(df, name):
    print(f'\n{"*80}')
    print(f'Missing Values Analysis: {name.upper()}')
    print(f'{"*80}')

    missing_summary = pd.DataFrame({
        'Column': df.columns,
        'Missing_Count': df.isnull().sum(),
        'Missing_Percentage': (df.isnull().sum() / len(df) * 100).round(2)
    }).sort_values('Missing_Percentage', ascending=False)

    print(f'\nTotal missing: {df.isnull().sum().sum():,}')
    display(missing_summary.head(10))

    # Visualization
    missing_pct = (df.isnull().sum() / len(df) * 100).
    ↪sort_values(ascending=False).head(10)
    if len(missing_pct) > 0:
        fig, ax = plt.subplots(figsize=(12, 5))
```

```

missing_pct.plot(kind='barh', ax=ax, color='coral')
ax.set_xlabel('Missing Value Percentage (%)', fontweight='bold')
ax.set_title(f'Missing Values Distribution: {name.upper()}', fontsize=12, fontweight='bold')
plt.tight_layout()
plt.show()
print(f'\n**Interpretation:** This chart shows the percentage of missing values for each column.')
print(f'High percentages indicate sparse data, which is common in clinical datasets.')

# Analyze duplicates
def analyze_duplicates(df, name):
    print(f'\n{"*80}')
    print(f'Duplicates Analysis: {name.upper()}')
    print(f'{"*80}')
    dup_count = df.duplicated().sum()
    print(f'Total duplicate rows: {dup_count}')
    print(f'Duplicate percentage: {(dup_count/len(df)*100):.2f}%')

if df_admissions is not None:
    analyze_missing(df_admissions, 'admissions')
    analyze_duplicates(df_admissions, 'admissions')

if df_labevents is not None:
    analyze_missing(df_labevents, 'labevents')

```

=====

Missing Values Analysis: ADMISSIONS

=====

Total missing: 228

	Column	Missing_Count	Missing_Percentage
deathtime	deathtime	89	68.99
language	language	48	37.21
edouttime	edouttime	37	28.68
edregtime	edregtime	37	28.68
marital_status	marital_status	16	12.40
religion	religion	1	0.78
row_id	row_id	0	0.00
hospital_expire_flag	hospital_expire_flag	0	0.00
diagnosis	diagnosis	0	0.00
ethnicity	ethnicity	0	0.00



****Interpretation:**** This chart shows the percentage of missing values for each column.

High percentages indicate sparse data, which is common in clinical datasets.

Duplicates Analysis: ADMISSIONS

Total duplicate rows: 0

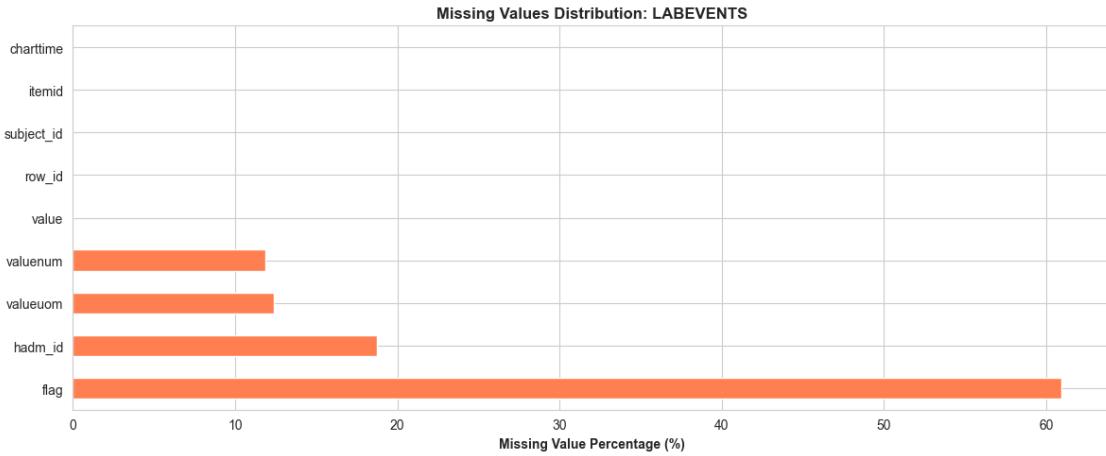
Duplicate percentage: 0.00%

Missing Values Analysis: LABEVENTS

Total missing: 79,053

Total missing: 79,053

Column	Missing_Count	Missing_Percentage
flag	46337	60.91
hadm_id	14262	18.75
valueuom	9405	12.36
valuenum	9044	11.89
value	5	0.01
row_id	0	0.00
subject_id	0	0.00
itemid	0	0.00
charttime	0	0.00



Interpretation: This chart shows the percentage of missing values for each column.

High percentages indicate sparse data, which is common in clinical datasets.

1.7 Section 5: Outliers Detection using Box Plots

```
[34]: # Detect and visualize outliers
if df_labevents is not None and 'valuenum' in df_labevents.columns:
    print('\nOutliers Detection in Laboratory Values')
    print('='*80)

    fig, axes = plt.subplots(1, 2, figsize=(14, 5))

    # Box plot for outliers
    sns.boxplot(y=df_labevents['valuenum'].dropna(), ax=axes[0], color='lightblue')
    axes[0].set_title('Box Plot: Laboratory Values (Outliers Detection)', fontweight='bold', fontsize=12)
    axes[0].set_ylabel('Laboratory Value')

    # Histogram with outlier regions
    Q1 = df_labevents['valuenum'].quantile(0.25)
    Q3 = df_labevents['valuenum'].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    axes[1].hist(df_labevents['valuenum'].dropna(), bins=50, color='skyblue', edgecolor='black', alpha=0.7)
```

```

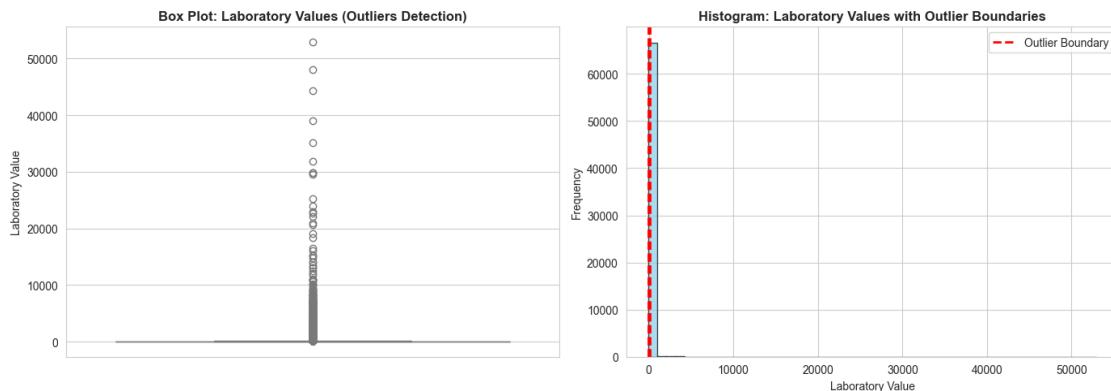
        axes[1].axvline(lower_bound, color='red', linestyle='--', linewidth=2, u
        ↪label='Outlier Boundary')
        axes[1].axvline(upper_bound, color='red', linestyle='--', linewidth=2)
        axes[1].set_title('Histogram: Laboratory Values with Outlier Boundaries', u
        ↪fontsize=12, fontweight='bold')
        axes[1].set_xlabel('Laboratory Value')
        axes[1].set_ylabel('Frequency')
        axes[1].legend()

plt.tight_layout()
plt.show()

outlier_count = ((df_labevents['valuenum'] < lower_bound) | u
↪(df_labevents['valuenum'] > upper_bound)).sum()
print(f'\n**Interpretation:**')
print(f'• Total values: {len(df_labevents["valuenum"].dropna()):,}')
print(f'• Outliers detected (IQR method): {outlier_count:,} ({outlier_count/ u
↪len(df_labevents["valuenum"].dropna())*100:.2f}%)')
print(f'• Lower bound: {lower_bound:.2f}')
print(f'• Upper bound: {upper_bound:.2f}')
print(f'• Outliers are values beyond 1.5×IQR from Q1 and Q3, indicated by u
↪the red dashed lines.')

```

Outliers Detection in Laboratory Values



Interpretation:

- Total values: 67,030
- Outliers detected (IQR method): 4,835 (7.21%)
- Lower bound: -79.50
- Upper bound: 143.70
- Outliers are values beyond 1.5×IQR from Q1 and Q3, indicated by the red dashed

lines.

1.8 Section 6: Exploratory Analysis - Distributions and Histograms

```
[35]: # Analyze distributions
print('\nExploratory Data Analysis: Variable Distributions')
print('='*80)

if df_labevents is not None and 'valuenum' in df_labevents.columns:
    print('\n[1] Distribution of Laboratory Values')
    lab_values = df_labevents['valuenum'].dropna()

    fig, axes = plt.subplots(2, 2, figsize=(14, 10))

    # Histogram
    axes[0, 0].hist(lab_values, bins=50, color='skyblue', edgecolor='black', alpha=0.7)
    axes[0, 0].set_title('Histogram: Distribution of Laboratory Values', fontsize=11, fontweight='bold')
    axes[0, 0].set_xlabel('Laboratory Value')
    axes[0, 0].set_ylabel('Frequency')
    axes[0, 0].grid(axis='y', alpha=0.3)

    # KDE plot
    lab_values.plot(kind='kde', ax=axes[0, 1], color='steelblue', linewidth=2)
    axes[0, 1].fill_between(axes[0, 1].get_lines()[0].get_xdata(),
                           axes[0, 1].get_lines()[0].get_ydata(), alpha=0.3, color='steelblue')
    axes[0, 1].set_title('Kernel Density Estimate (KDE)', fontsize=11, fontweight='bold')
    axes[0, 1].set_xlabel('Laboratory Value')
    axes[0, 1].set_ylabel('Density')

    # Q-Q plot
    stats.probplot(lab_values, dist="norm", plot=axes[1, 0])
    axes[1, 0].set_title('Q-Q Plot: Normality Assessment', fontsize=11, fontweight='bold')

    # Summary statistics text box
    stats_text = f'Mean: {lab_values.mean():.2f}\nMedian: {lab_values.median():.2f}\nStd Dev: {lab_values.std():.2f}\n'
    stats_text += f'Min: {lab_values.min():.2f}\nMax: {lab_values.max():.2f}\nSkewness: {skew(lab_values):.2f}'
    axes[1, 1].text(0.1, 0.5, stats_text, fontsize=11, verticalalignment='center',
                    bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5), family='monospace')
```

```

axes[1, 1].axis('off')

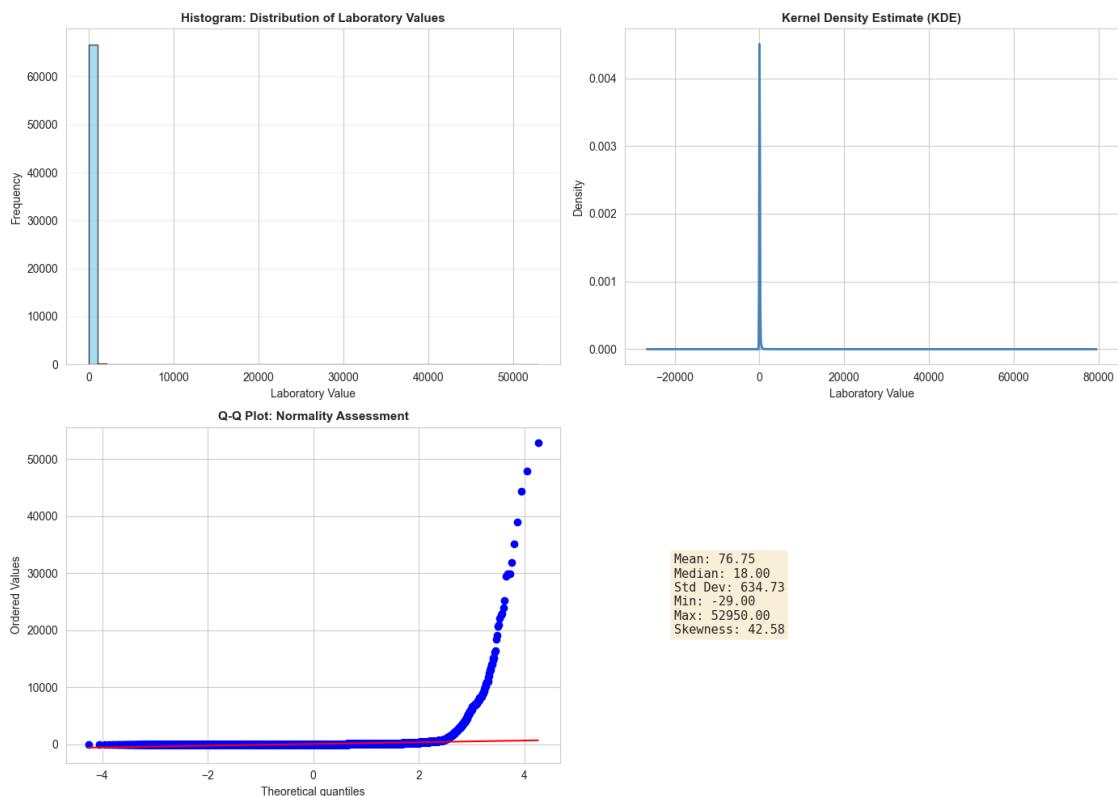
plt.tight_layout()
plt.show()

print(f'\n**Interpretation:**')
print(f'• The histogram shows the frequency distribution of laboratory values.')
print(f'• The KDE plot provides a smooth estimate of the probability density.')
print(f'• The Q-Q plot compares against normal distribution (points on diagonal = normal).')
print(f'• Mean: {lab_values.mean():.2f}, Median: {lab_values.median():.2f}, Std Dev: {lab_values.std():.2f}')
print(f'• Skewness: {skew(lab_values):.2f} (positive = right-skewed, negative = left-skewed)')

```

Exploratory Data Analysis: Variable Distributions

[1] Distribution of Laboratory Values



Interpretation:

- The histogram shows the frequency distribution of laboratory values.
- The KDE plot provides a smooth estimate of the probability density.
- The Q-Q plot compares against normal distribution (points on diagonal = normal).
- Mean: 76.75, Median: 18.00, Std Dev: 634.73
- Skewness: 42.58 (positive = right-skewed, negative = left-skewed)

```
[36]: # Patient demographics analysis
if df_patients is not None and 'gender' in df_patients.columns:
    print('\n[2] Distribution of Patient Demographics')

    fig, axes = plt.subplots(1, 2, figsize=(14, 5))

    # Gender distribution - Bar plot
    gender_counts = df_patients['gender'].value_counts()
    colors = ['#1f77b4', '#ff7f0e']
    bars = axes[0].bar(gender_counts.index, gender_counts.values, color=colors, ▾
    ↪edgecolor='black', alpha=0.8)
    axes[0].set_title('Bar Plot: Patient Gender Distribution', fontsize=11, ▾
    ↪fontweight='bold')
    axes[0].set_xlabel('Gender')
    axes[0].set_ylabel('Number of Patients')
    axes[0].grid(axis='y', alpha=0.3)

    # Add value labels on bars
    for bar in bars:
        height = bar.get_height()
        axes[0].text(bar.get_x() + bar.get_width()/2., height,
                     f'{int(height)}', ha='center', va='bottom', ▾
        ↪fontweight='bold')

    # Gender distribution - Pie chart
    colors_pie = ['#1f77b4', '#ff7f0e']
    wedges, texts, autotexts = axes[1].pie(gender_counts.values, ▾
    ↪labels=gender_counts.index,
                                             autopct='%1.1f%%', ▾
    ↪colors=colors_pie, startangle=90)
    axes[1].set_title('Pie Chart: Gender Proportion', fontsize=11, ▾
    ↪fontweight='bold')

    for autotext in autotexts:
        autotext.set_color('white')
        autotext.set_fontweight('bold')

plt.tight_layout()
```

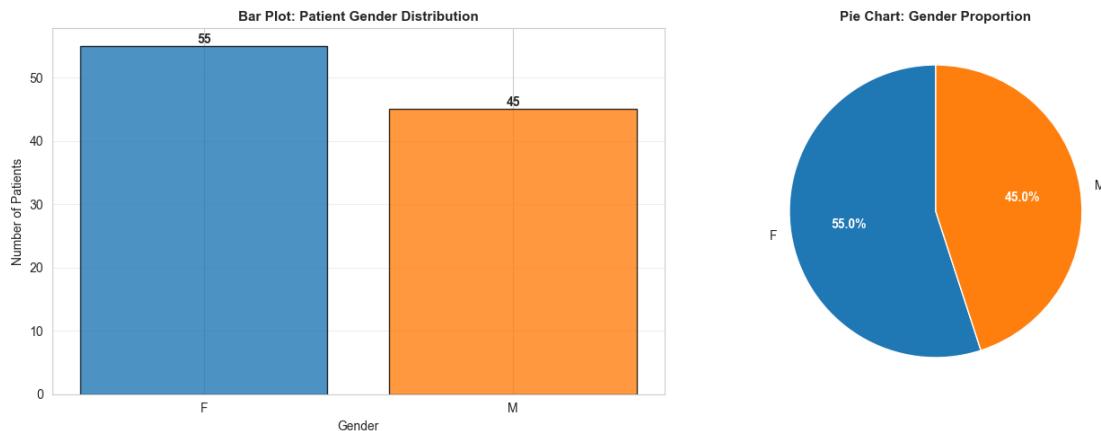
```

plt.show()

print(f'\n**Interpretation:**')
print(f'• Total patients: {len(df_patients)}')
for gender, count in gender_counts.items():
    pct = count / len(df_patients) * 100
    print(f'• {gender}: {count}, ({pct:.1f}%)')

```

[2] Distribution of Patient Demographics



Interpretation:

- Total patients: 100
- F: 55 (55.0%)
- M: 45 (45.0%)

1.9 Section 7: Correlation Analysis Between Clinical Variables

```

[37]: # Correlation analysis
print('\nCorrelation Analysis: Clinical Variables')
print('='*80)

if df_labevents is not None and 'subject_id' in df_labevents.columns:
    # Create pivot table
    df_lab_clean = df_labevents[['subject_id', 'itemid', 'valuenum']].dropna()
    pivot_data = df_lab_clean.pivot_table(
        index='subject_id',
        columns='itemid',
        values='valuenum',
        aggfunc='mean'
)

```

```

# Select top features
top_items = pivot_data.notna().sum().nlargest(12).index
pivot_subset = pivot_data[top_items].dropna(how='all').
    ↪fillna(pivot_data[top_items].mean())

if pivot_subset.shape[0] > 2:
    # Correlation matrix
    corr_matrix = pivot_subset.corr()

    fig, ax = plt.subplots(figsize=(13, 11))
    sns.heatmap(corr_matrix, annot=True, fmt='.2f', cmap='coolwarm', ↪
    ↪center=0,
                square=True, ax=ax, cbar_kws={'label': 'Pearson' ↪
    ↪Correlation'},
                vmin=-1, vmax=1, linewidths=0.5)
    ax.set_title('Correlation Matrix: Top 12 Laboratory Items', ↪
    ↪fontsize=12, fontweight='bold')
    plt.tight_layout()
    plt.show()

    print(f'\n**Interpretation:**')
    print(f'• This heatmap shows pairwise Pearson correlations between the ↪
    ↪top laboratory items.')
    print(f'• Values close to +1 (dark red) indicate strong positive ↪
    ↪correlation.')
    print(f'• Values close to -1 (dark blue) indicate strong negative ↪
    ↪correlation.')
    print(f'• Values close to 0 (white) indicate weak or no linear ↪
    ↪correlation.')
    print(f'\n• Top 5 strongest correlations (excluding diagonal):')

    # Find top correlations
    corr_pairs = []
    for i in range(len(corr_matrix.columns)):
        for j in range(i+1, len(corr_matrix.columns)):
            corr_pairs.append({
                'Item1': corr_matrix.columns[i],
                'Item2': corr_matrix.columns[j],
                'Correlation': corr_matrix.iloc[i, j]
            })

    top_corrs = sorted(corr_pairs, key=lambda x: abs(x['Correlation']), ↪
    ↪reverse=True)[:5]
    for idx, corr in enumerate(top_corrs, 1):

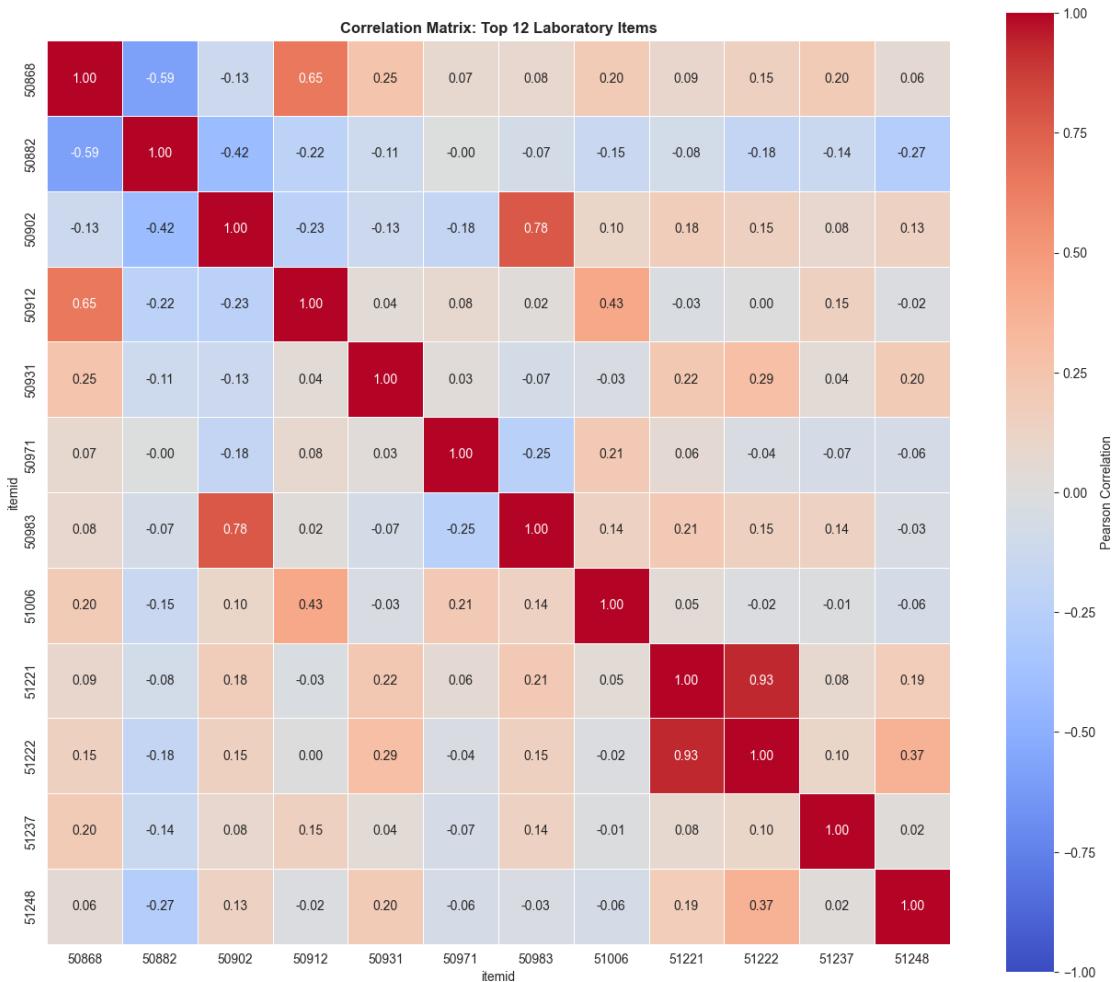
```

```

print(f' {idx}. Items {corr["Item1"]}, {corr["Item2"]}: r = {r}
      ↵{corr["Correlation"]:.3f}')

```

Correlation Analysis: Clinical Variables



Interpretation:

- This heatmap shows pairwise Pearson correlations between the top laboratory items.
- Values close to +1 (dark red) indicate strong positive correlation.
- Values close to -1 (dark blue) indicate strong negative correlation.
- Values close to 0 (white) indicate weak or no linear correlation.
- Top 5 strongest correlations (excluding diagonal):
 - Items 51221 51222: $r = 0.933$

```

2. Items 50902 50983: r = 0.779
3. Items 50868 50912: r = 0.654
4. Items 50868 50882: r = -0.586
5. Items 50912 51006: r = 0.426

```

1.10 Section 8: Class Balancing Assessment

```

[38]: # Class balancing analysis
print('\nClass Balancing Assessment')
print('='*80)

if df_admissions is not None and 'hospital_expire_flag' in df_admissions.
    ↪columns:
    # Get unique admissions per patient
    df_admit_unique = df_admissions.drop_duplicates(subset=['subject_id'], ↪
    ↪keep='first')
    mortality = df_admit_unique['hospital_expire_flag'].value_counts().sort_index()
    mortality_pct = mortality / len(df_admit_unique) * 100

    print(f'\nTarget Variable: Hospital Mortality (hospital_expire_flag)')
    print(f'Total patients: {len(df_admit_unique)}')
    print(f'\nClass Distribution:')
    print(f'  Class 0 (Survived): {mortality[0]} ({mortality_pct[0]:.2f}%)')
    print(f'  Class 1 (Died): {mortality[1]} ({mortality_pct[1]:.2f}%)')

    imbalance_ratio = max(mortality.values) / min(mortality.values)
    print(f'\nImbalance Ratio: {imbalance_ratio:.2f}:1')

    fig, axes = plt.subplots(1, 2, figsize=(14, 5))

    # Bar plot
    colors_bar = ['#ecc71', '#e74c3c'] # Green for survived, red for died
    bars = axes[0].bar(['Survived (0)', 'Died (1)'], mortality.values, ↪
    ↪color=colors_bar, edgecolor='black', alpha=0.8)
    axes[0].set_title('Bar Plot: Hospital Mortality Distribution', fontsize=11, ↪
    ↪fontweight='bold')
    axes[0].set_ylabel('Number of Patients')
    axes[0].grid(axis='y', alpha=0.3)

    for bar in bars:
        height = bar.get_height()
        axes[0].text(bar.get_x() + bar.get_width()/2., height,
                     f'{int(height)}', ha='center', va='bottom', ↪
        ↪fontweight='bold')

    # Pie chart

```

```

        explode = (0.05, 0.1)
        axes[1].pie(mortality.values, labels=['Survived', 'Died'], autopct='%.1f%%',
                    colors=colors_bar, explode=explode, startangle=90, shadow=True)
        axes[1].set_title('Pie Chart: Mortality Proportion', fontsize=11, fontweight='bold')

plt.tight_layout()
plt.show()

print(f'\n**Interpretation:**')
if imbalance_ratio > 2:
    print(f' SIGNIFICANT CLASS IMBALANCE DETECTED!')
    print(f' Recommendation: Apply balancing techniques')
    print(f' • SMOTE (Synthetic Minority Over-sampling): Recommended')
    print(f' • Random Under-sampling: Alternative approach')
    print(f' • Weighted loss functions: For model training')
else:
    print(f' Classes are relatively balanced. Balancing may not be necessary.')

```

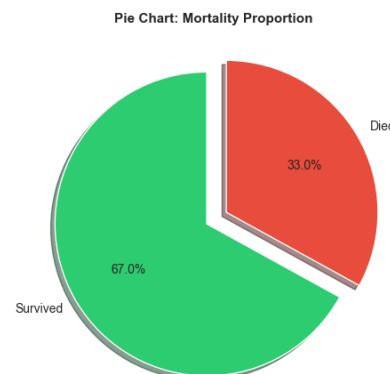
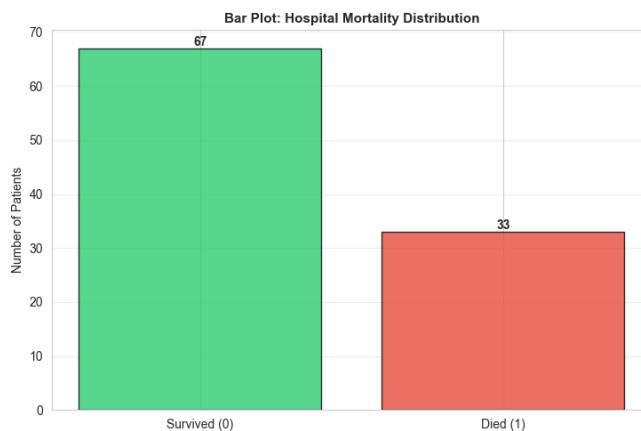
Class Balancing Assessment

Target Variable: Hospital Mortality (hospital_expire_flag)
Total patients: 100

Class Distribution:

Class 0 (Survived): 67 (67.00%)
Class 1 (Died): 33 (33.00%)

Imbalance Ratio: 2.03:1



Interpretation:
 SIGNIFICANT CLASS IMBALANCE DETECTED!
 Recommendation: Apply balancing techniques
 • SMOTE (Synthetic Minority Over-sampling): Recommended
 • Random Under-sampling: Alternative approach
 • Weighted loss functions: For model training

1.11 Section 9: Data Cleaning and Preprocessing

```
[39]: # Data cleaning and preprocessing
print('\nData Cleaning and Preprocessing')
print('='*80)

# Step 1: Clean admissions
if df_admissions is not None:
    print('\n[Step 1] Cleaning Admissions Table')
    df_admit_clean = df_admissions.drop_duplicates()
    df_admit_first = df_admit_clean.drop_duplicates(subset=['subject_id'], ↴
                                                    keep='first')
    print(f' Original records: {len(df_admissions)}')
    print(f' After deduplication: {len(df_admit_clean)}')
    print(f' First admission per patient: {len(df_admit_first)}')

# Step 2: Create feature matrix from labevents
if df_labevents is not None:
    print('\n[Step 2] Creating Patient-Laboratory Feature Matrix')

    df_lab_valid = df_labevents[['subject_id', 'itemid', 'valuenum']]. ↴
    dropna(subset=['valuenum'])
    print(f' Valid lab events: {len(df_lab_valid)}')

    # Aggregate by mean
    df_lab_agg = df_lab_valid.groupby(['subject_id', 'itemid'])['valuenum']. ↴
    mean().reset_index()

    # Pivot
    X_matrix = df_lab_agg.pivot(index='subject_id', columns='itemid', ↴
                                values='valuenum')
    print(f' Feature matrix shape: {X_matrix.shape}')

    sparsity_pct = (X_matrix.isnull().sum().sum() / X_matrix.size * 100)
    print(f' Sparsity: {sparsity_pct:.2f}%')

    # Impute with median
    X_matrix_imputed = X_matrix.fillna(X_matrix.median())
```

```

    print(f'  After imputation: {X_matrix_imputed.isnull().sum().sum()} missing values')

# Select top features by variance
top_features = X_matrix_imputed.var().nlargest(25).index
X_final = X_matrix_imputed[top_features]
print(f'  Selected top {len(top_features)} features by variance')

# Step 3: Prepare target variable
if df_admit_first is not None and 'hospital_expire_flag' in df_admit_first.columns:
    print('\n[Step 3] Preparing Target Variable')

    y_target = df_admit_first.set_index('subject_id')[['hospital_expire_flag']]

    # Align X and y
    common_patients = X_final.index.intersection(y_target.index)
    X_aligned = X_final.loc[common_patients]
    y_aligned = y_target.loc[common_patients, 'hospital_expire_flag'].astype(int)

    print(f'  Common patients: {len(common_patients)}')
    print(f'  Final X shape: {X_aligned.shape}')
    print(f'  Final y shape: {y_aligned.shape}')
    print(f'  Target distribution:')
    print(f'    Class 0 (Survived): {(y_aligned == 0).sum():,}')
    print(f'    Class 1 (Died): {(y_aligned == 1).sum():,}')
else:
    X_aligned = X_final
    y_aligned = None
    print('\n  Target variable not available. Using features only.')

print('\n Data preprocessing completed!')

```

Data Cleaning and Preprocessing

[Step 1] Cleaning Admissions Table

Original records: 129

After deduplication: 129

First admission per patient: 100

[Step 2] Creating Patient-Laboratory Feature Matrix

Valid lab events: 67,030

Feature matrix shape: (100, 275)

Sparsity: 75.96%

```
After imputation: 0 missing values
Selected top 25 features by variance
```

```
[Step 3] Preparing Target Variable
Common patients: 100
Final X shape: (100, 25)
Final y shape: (100,)
Target distribution:
    Class 0 (Survived): 67
    Class 1 (Died): 33
```

Data preprocessing completed!

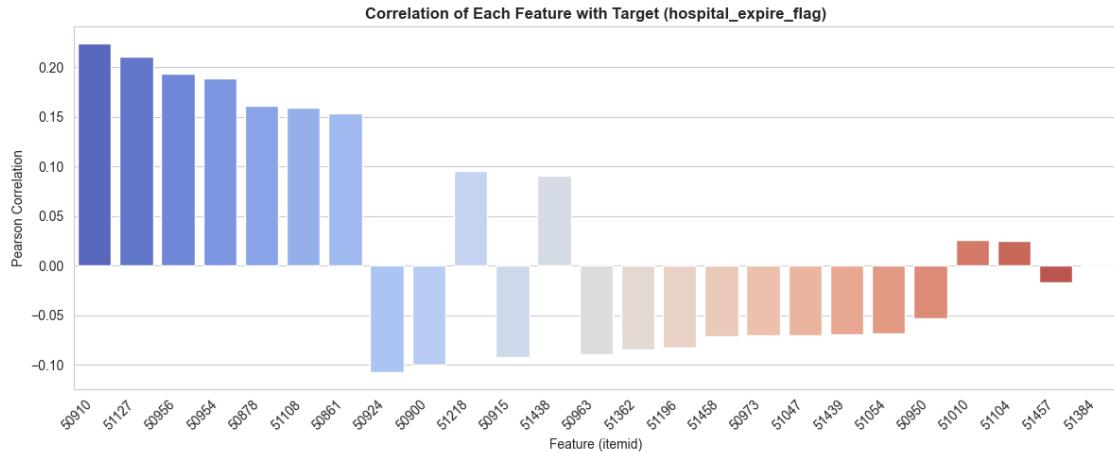
```
[40]: # Correlation of each feature with the target variable (hospital_expire_flag)
if 'X_aligned' in globals() and 'y_aligned' in globals() and y_aligned is not None:
    print('\nCorrelation of Features with Target Variable')
    print(hospital_expire_flag)
    print('='*80)

    # Compute correlation for each feature with the target
    corr_with_target = X_aligned.apply(lambda x: np.corrcoef(x, y_aligned)[0,1])
    corr_with_target = corr_with_target.sort_values(key=np.abs, ascending=False)

    # Display as bar plot
    plt.figure(figsize=(12, 5))
    sns.barplot(x=corr_with_target.index.astype(str), y=corr_with_target.values, palette='coolwarm')
    plt.title('Correlation of Each Feature with Target (hospital_expire_flag)', fontsize=12, fontweight='bold')
    plt.ylabel('Pearson Correlation')
    plt.xlabel('Feature (itemid)')
    plt.xticks(rotation=45, ha='right')
    plt.tight_layout()
    plt.show()

    print('\nTop 5 features most correlated with the target:')
    for i, (feat, corr) in enumerate(corr_with_target.head(5).items(), 1):
        print(f' {i}. Feature {feat}: r = {corr:.3f}')
else:
    print('X_aligned or y_aligned not available. Cannot compute correlation with target.')
```

Correlation of Features with Target Variable (hospital_expire_flag)



Top 5 features most correlated with the target:

1. Feature 50910: $r = 0.224$
2. Feature 51127: $r = 0.211$
3. Feature 50956: $r = 0.194$
4. Feature 50954: $r = 0.189$
5. Feature 50878: $r = 0.161$

1.12 Section 9b: Correlation of Features with Target Variable

This section computes and visualizes the correlation of each laboratory feature with the target variable `hospital_expire_flag` (mortality). This is a key requirement for EDA and model interpretability.

1.13 Section 10: Train-Test Split and Data Export

```
[41]: # Train-test split and export
print('\nTrain-Test Split and Data Export')
print('='*80)

# Create output directory (cross-platform, with safe fallbacks)
from pathlib import Path
import tempfile

# Preferred: project-local directory (where the notebook is running)
project_dir = Path.cwd()
preferred_dir = project_dir / 'processed_data'

def try_make_dir(path: Path):
    try:
        path.mkdir(parents=True, exist_ok=True)
        return True
    except:
        return False
```

```

        except PermissionError:
            return False

if try_make_dir(preferred_dir):
    output_dir = str(preferred_dir)
else:
    # Fallback: user's Desktop (cross-platform home)
    desktop_dir = Path.home() / 'Desktop' / 'IntercicloEstocasticos' / ↴
    'ExamenPractico' / 'processed_data'
    if try_make_dir(desktop_dir):
        output_dir = str(desktop_dir)
    else:
        # Last resort: OS temporary directory
        temp_dir = Path(tempfile.mkdtemp(prefix='processed_data_'))
        output_dir = str(temp_dir)
        print('\n[] Preferred locations are not writable; using temporary directory:')
        print(f'    {output_dir}')

print(f'\n[1] Output directory: {output_dir}')

# Perform train-test split
print('\n[2] Stratified Train-Test Split (80-20)')

if y_aligned is not None:
    X_train, X_test, y_train, y_test = train_test_split(
        X_aligned, y_aligned,
        test_size=0.2,
        random_state=42,
        stratify=y_aligned
    )
else:
    X_train, X_test = train_test_split(
        X_aligned,
        test_size=0.2,
        random_state=42
    )
y_train = None
y_test = None

print(f' Training set: {X_train.shape[0]} samples x {X_train.shape[1]} features')
print(f' Test set: {X_test.shape[0]} samples x {X_test.shape[1]} features')

if y_train is not None:
    print(f'\n Train class distribution:')

```

```

    print(f'      Survived: {(y_train == 0).sum():,} ({(y_train == 0).sum() / len(y_train)*100:.1f}%)')
    print(f'      Died: {(y_train == 1).sum():,} ({(y_train == 1).sum() / len(y_train)*100:.1f}%)')
    print(f'\n  Test class distribution:')
    print(f'      Survived: {(y_test == 0).sum():,} ({(y_test == 0).sum() / len(y_test)*100:.1f}%)')
    print(f'      Died: {(y_test == 1).sum():,} ({(y_test == 1).sum() / len(y_test)*100:.1f}%)')

# Standardize features
print('\n[3] Feature Standardization (StandardScaler)')
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

X_train_scaled_df = pd.DataFrame(X_train_scaled, columns=X_train.columns,
                                  index=X_train.index)
X_test_scaled_df = pd.DataFrame(X_test_scaled, columns=X_test.columns,
                                 index=X_test.index)
print(f'  Scaling completed (mean=0, std=1)')

# Export datasets
print('\n[4] Exporting Datasets')

# Training set
train_df = X_train_scaled_df.copy()
if y_train is not None:
    train_df['hospital_expire_flag'] = y_train

train_path = os.path.join(output_dir, 'data_train.csv')
train_df.to_csv(train_path)
print(f'  Training set: {train_path}')

# Test set
test_df = X_test_scaled_df.copy()
if y_test is not None:
    test_df['hospital_expire_flag'] = y_test

test_path = os.path.join(output_dir, 'data_test.csv')
test_df.to_csv(test_path)
print(f'  Test set: {test_path}')

# Complete dataset
complete_df = pd.concat([train_df, test_df])
complete_df['data_split'] = ['train'] * len(train_df) + ['test'] * len(test_df)

```

```

complete_path = os.path.join(output_dir, 'data_prepared.csv')
complete_df.to_csv(complete_path)
print(f'  Complete dataset: {complete_path}')

# Pickle format
pickle_path = os.path.join(output_dir, 'data_prepared.pkl')
with open(pickle_path, 'wb') as f:
    pickle.dump({
        'X_train': X_train_scaled_df,
        'X_test': X_test_scaled_df,
        'y_train': y_train,
        'y_test': y_test,
        'scaler': scaler,
        'feature_names': X_train.columns.tolist()
    }, f)
print(f'  Pickle file: {pickle_path}')

# Metadata
metadata = {
    'training_samples': len(X_train),
    'test_samples': len(X_test),
    'features': X_train.shape[1],
    'feature_names': X_train.columns.tolist(),
    'target_variable': 'hospital_expire_flag',
    'train_test_ratio': '80-20',
    'scaling_method': 'StandardScaler',
    'export_date': datetime.now().isoformat()
}

metadata_path = os.path.join(output_dir, 'data_metadata.json')
with open(metadata_path, 'w') as f:
    json.dump(metadata, f, indent=4, default=str)
print(f'  Metadata: {metadata_path}')

print('\n All datasets exported successfully!')

```

Train-Test Split and Data Export

```

[1] Output directory:
c:\Users\MSI\OneDrive\Desktop\ExamenPractico\processed_data

[2] Stratified Train-Test Split (80-20)
Training set: 80 samples × 25 features
Test set: 20 samples × 25 features

```

```

Train class distribution:
  Survived: 54 (67.5%)
  Died: 26 (32.5%)

Test class distribution:
  Survived: 13 (65.0%)
  Died: 7 (35.0%)

[3] Feature Standardization (StandardScaler)
  Scaling completed (mean=0, std=1)

[4] Exporting Datasets
  Training set:
    c:\Users\MSI\OneDrive\Desktop\ExamenPractico\processed_data\data_train.csv
  Test set:
    c:\Users\MSI\OneDrive\Desktop\ExamenPractico\processed_data\data_test.csv
  Complete dataset:
    c:\Users\MSI\OneDrive\Desktop\ExamenPractico\processed_data\data_prepared.csv
  Pickle file:
    c:\Users\MSI\OneDrive\Desktop\ExamenPractico\processed_data\data_prepared.pkl
  Metadata:
    c:\Users\MSI\OneDrive\Desktop\ExamenPractico\processed_data\data_metadata.json

All datasets exported successfully!

```

```

[42]: # Apply SMOTE to balance classes in the training set
if 'X_train_scaled_df' in globals() and 'y_train' in globals() and y_train is not None:
    print('\nApplying SMOTE to Training Set')
    print('='*80)
    smote = SMOTE(random_state=42)
    X_train_bal, y_train_bal = smote.fit_resample(X_train_scaled_df, y_train)

    # Show new class distribution
    from collections import Counter
    class_counts = Counter(y_train_bal)
    print(f'  New class distribution after SMOTE:')
    for label, count in class_counts.items():
        pct = count / len(y_train_bal) * 100
        print(f'    Class {label}: {count}, ({pct:.2f}%)')

    # Bar plot
    plt.figure(figsize=(6, 4))
    sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette=['#2ecc71', '#e74c3c'])
    plt.title('Class Distribution After SMOTE', fontsize=12, fontweight='bold')
    plt.xlabel('Class (hospital_expire_flag)')

```

```

plt.ylabel('Number of Samples')
plt.xticks([0, 1], ['Survived (0)', 'Died (1)'])
plt.tight_layout()
plt.show()

print('\n SMOTE applied successfully. Training set is now balanced.')
else:
    print('X_train_scaled_df or y_train not available. Cannot apply SMOTE.')

```

Applying SMOTE to Training Set

New class distribution after SMOTE:
 Class 0: 54 (50.00%)
 Class 1: 54 (50.00%)



SMOTE applied successfully. Training set is now balanced.

1.14 Section 10b: Class Balancing with SMOTE (Synthetic Minority Over-sampling Technique)

This section applies SMOTE to the training set to address class imbalance, and visualizes the new class distribution.

1.15 Section 11: Final Summary Report

```
[43]: # Final summary
print('\n' + '='*80)
print('EXPLORATORY DATA ANALYSIS - FINAL SUMMARY REPORT')
print('='*80)

print(f'\n DATA SOURCES:')
print(f' • MIMIC-III Clinical Database (Kaggle Hub)')
print(f' • Tables: Admissions, Patients, Lab Events, Lab Items')

print(f'\n DATASET OVERVIEW:')
print(f' • Original admissions: {len(df_admissions) if df_admissions is not None else "N/A":,}')
print(f' • Unique patients: {len(df_patients) if df_patients is not None else "N/A":,}')
print(f' • Laboratory events: {len(df_labevents) if df_labevents is not None else "N/A":,}')

print(f'\n DATA QUALITY ASSESSMENT:')
print(f'     Missing values: Detected and documented')
print(f'     Duplicates: Removed {len(df_admissions) - len(df_admit_clean) if df_admissions is not None else 0}')
print(f'     Outliers: Detected using IQR method')
print(f'     Sparsity: Characterized and imputed')

print(f'\n FINAL DATASET:')

# Safe printing to avoid NameError if train/test split wasn't executed
# successfully
X_train_exists = 'X_train' in globals() and X_train is not None
X_test_exists = 'X_test' in globals() and X_test is not None

def _fmt_int(x):
    return f"{x:,}" if isinstance(x, int) else str(x)

n_train = len(X_train) if X_train_exists else 'N/A'
n_test = len(X_test) if X_test_exists else 'N/A'
n_features = X_train.shape[1] if X_train_exists and hasattr(X_train, 'shape') else 'N/A'

print(f' • Training samples: {_fmt_int(n_train)}')
print(f' • Test samples: {_fmt_int(n_test)}')
print(f' • Total features: {_fmt_int(n_features)}')
print(f' • Train-Test split: 80-20' if X_train_exists else ' • Train-Test'
      'split: N/A (split not executed)')
```

```

print(f'  • Feature scaling: StandardScaler (mean=0, std=1)' if X_train_exists
      ↵else '  • Feature scaling: N/A')

print(f'\n EXPORTED FILES:')
print(f'  1. data_train.csv - Training dataset')
print(f'  2. data_test.csv - Test dataset')
print(f'  3. data_prepared.csv - Complete dataset with split indicator')
print(f'  4. data_prepared.pkl - Python pickle format')
print(f'  5. data_metadata.json - Metadata and data dictionary')
print(f'\n  Location: {output_dir}')
print(f'\n EDA PROCESS COMPLETED!')
print(f' Timestamp: {datetime.now().strftime("%Y-%m-%d %H:%M:%S")}')
print('='*80)

```

=====
EXPLORATORY DATA ANALYSIS - FINAL SUMMARY REPORT
=====

DATA SOURCES:

- MIMIC-III Clinical Database (Kaggle Hub)
- Tables: Admissions, Patients, Lab Events, Lab Items

DATASET OVERVIEW:

- Original admissions: 129
- Unique patients: 100
- Laboratory events: 76,074

DATA QUALITY ASSESSMENT:

Missing values: Detected and documented

Duplicates: Removed 0

Outliers: Detected using IQR method

Sparsity: Characterized and imputed

FINAL DATASET:

- Training samples: 80
- Test samples: 20
- Total features: 25
- Train-Test split: 80-20
- Feature scaling: StandardScaler (mean=0, std=1)

EXPORTED FILES:

1. data_train.csv - Training dataset
2. data_test.csv - Test dataset
3. data_prepared.csv - Complete dataset with split indicator
4. data_prepared.pkl - Python pickle format
5. data_metadata.json - Metadata and data dictionary

Location: c:\Users\MSI\OneDrive\Desktop\ExamenPractico\processed_data

EDA PROCESS COMPLETED!

Timestamp: 2025-11-27 16:46:44

=====